

# Food Delivery Project using Spring

The objective of this project is to develop a food delivery application (similar to Zomato, Swiggy). The application will be organized as a set of three microservices, each hosting a RESTful service. The high-level features of the application are as follows:

- There is a fixed set of [customers](#), [restaurants](#), [item IDs](#), and [delivery agents](#), specified initially in a given input file when the application starts. This input file is called `/initialData.txt`, and is explained more in a section titled Initialization later in this document.
- The three services are [Restaurant](#), [Delivery](#), and [Wallet](#).
- Delivery is the main service, with which customers interact and which invokes the other services in turn. The Delivery service keeps track of orders placed so far, the current statuses of the orders, and the current statuses of the delivery agents. The Delivery service is also aware of the price of each item in each restaurant (this is constant and specified upfront). This way, the Delivery service is able to calculate the total amount of an order when the order is received. An order can be for a single item only (but any quantity of it).
- The Restaurant service keeps track of the inventory of items available in all the restaurants. The inventory reduces when an order is received, and can be increased using a specified end-point.
- The Wallet service keeps track of the balance maintained by each customer, and supports end-points to decrease or increase wallet amounts.
- Any delivery agent can sign-in whenever they like. Whenever they are in signed-in state, they are either *available* (ready to deliver an order) or *unavailable* (i.e., currently delivering an order). Upon sign-in they start off in *available* state. They can sign-out whenever they like provided they are in *available* state. The Delivery service can assign an order to an agent only if the agent is in *available* state.
- Restaurants are assumed to be always open and serving.
- For any end-point below, if it is not mentioned who is to invoke it, then it is intended to be invoked by a human (delivery agent, customer, restaurant manager, etc). In our setting, a test-script will send requests to these end-points on behalf of humans.

## End-Points in [Restaurant](#) Service

1. POST `/acceptOrder`  
with JSON payload of the form `{"restId": num, "itemId": x, "qty": y}`

This end-point will be invoked by the Delivery service.

If the restaurant with ID *num* has at least *y* quantities of itemId *x* in its current inventory

then

reduce the inventory of item *x* in restaurant *num* by *y*

```
        return HTTP status code 201 (created)
    else
        return HTTP status code 410 (gone)
```

2. POST /refillItem  
with JSON payload of the form {"restId": num, "itemId": x, "qty": y}

Increases the inventory of itemId x in restaurant num by y (provided num supplies x as per /initialData.txt). Should return status code 201 unconditionally.

3. POST /reInitialize

Set inventory of all items in all restaurants as given in the /initialData.txt file.  
Return HTTP status code 201.

## End-points in Delivery Service

Keeps track of the status of each order placed so far. The status of any order can be either *unassigned*, *assigned*, or *delivered* (*unassigned* means not yet given to a delivery agent). For any *assigned* or *delivered* order, the service also records the agentId to which this order is assigned. The service also keeps track of the current status of each delivery agent (*signed-out*, *unavailable*, *available*). Maintains information about the price of each itemId in each restaurant (this is constant information).

1. POST /requestOrder  
with JSON payload of the form {"custId": num, "restId": x, "itemId": y, "qty": z}

This is the main end-point used by the customer to place an order. It first computes the total billing amount for this order (as it maintains the price of each itemId in each restaurant). It tries to deduct the amount from the custId num's wallet. If deduction does not succeed, return HTTP status code 410. If the deduction succeeds, invoke the Restaurant Service's acceptOrder end-point, passing to it x, y, z. If that returns 410, restore the wallet balance and return HTTP status code 410. Otherwise, first generate a fresh orderId w (let orderIds start from 1000 and let them be incremented by one each time). Initialize the status of this fresh order w as *unassigned*. Then see if any Delivery Agent is *available* right now. If yes, assign an agentId to this order w from among the available agentIds, update the status of the orderId w to *assigned*, mark the assigned agentId as *unavailable*, and record that w is assigned to this agentId. (If many agents are available, choose the lowest numbered agentId among them.) Then, whether or not the orderId was assigned an agent, return HTTP status code 201 and return response body {"orderId": w}.

2. POST /agentSignIn

with JSON payload of the form {"agentId": *num*}, where *num* is an agentId

If *num* is already in *available* or *unavailable* state, do nothing. Otherwise, If any orderIds are currently in *unassigned* state, find the least numbered orderId *y* that is unassigned, mark *num* as *unavailable*, mark the status of *y* as *assigned*, and record that *y* is assigned to *num*. Otherwise, mark the status of *num* as *available*. In all cases return HTTP status code 201.

3. POST /agentSignOut

with JSON payload of the form {"agentId": *num*}, where *num* is an agentId

If *num* is already *signed-out* or is *unavailable*, do nothing, else mark *num* as being in *signed-out* state. In both cases, return HTTP status code 201.

4. POST /orderDelivered

with JSON payload of the form {"orderId": *num*}, where *num* is an orderId

Ignore the request if *num* is not an orderId in *assigned* state. Otherwise, mark the status of orderId *num* as *delivered*, and mark the status of the agentId who was assigned this order as *available*. Let *x* be this agentId. If any orderIds are currently in *unassigned* state, find the least numbered orderId *y* that is unassigned, mark *x* as *unavailable* again, mark *y* as *assigned*, and record that *y* is assigned to *x*. In all cases return 201.

5. GET /order/*num*

where *num* is an orderId.

If *num* is a non-existent orderId return HTTP status code 404. Otherwise return status code 200 along with response JSON of the form {"orderId": *num*, "status": *x*, "agentId": *y*}, where *x* is *unassigned*, or *assigned*, or *delivered*. *y* will be the agentId that is assigned the order *num* in case *num* is in *assigned* or *delivered* state, else *y* will be -1.

6. GET /agent/*num*

where *num* is an agentId

Return status code 200 and response JSON of the form {"agentId": *num*, "status": *y*}, where *y* is *signed-out*, *available*, or *unavailable*.

7. POST /reInitialize

Delete all orders from the records (no matter what their status is), and mark status of each agent as *signed-out*. Return HTTP status code 201.

## End-points in **Wallet** service

1. POST /addBalance

with JSON payload of the form {"custId": *num*, "amount": *z*}

To be invoked by Delivery Service.

Increase the balance of custId *num* by *z*, and return HTTP status code 201.

2. POST /deductBalance

with JSON payload of the form {"custId": *num*, "amount": *z*}

To be invoked by Delivery Service.

If current balance of custId *num* is less than *z*, return HTTP status code 410, else reduce custId *num*'s balance by *z* and return HTTP status code 201.

3. GET /balance/*num*

where *num* is a custId. Return HTTP status code 200, and response JSON of the form {"custId": *num*, "balance": *z*}, where *z* is the current balance of custId *num*.

4. POST /reInitialize

Set balance of all customers to the initial value as given in the /initialData.txt file. Return HTTP status code 201.

## Requirements from implementation

You should develop each of the three services as a separate Spring project. Each project should be buildable using `./mvnw package`. Each project should contain a Dockerfile at the root folder. The Docker image from each project needs to be run (thus resulting in three separate microservices) in order to use the system. Using port forwarding in `docker run` you can set it up in a way that the Restaurant, Delivery, and Wallet services are available, respectively, at ports 8080, 8081, and 8082, on <http://localhost>.

When we run your project for evaluation purpose, we will use commands as shown here:

[https://indianinstituteofscience-my.sharepoint.com/:t:/g/personal/raghavan\\_iisc\\_ac\\_in/Eee9OQY3WNZKt4EtlvdgeccBvaiT5hMEc5hnbT1RD4545w?e=lb2fft](https://indianinstituteofscience-my.sharepoint.com/:t:/g/personal/raghavan_iisc_ac_in/Eee9OQY3WNZKt4EtlvdgeccBvaiT5hMEc5hnbT1RD4545w?e=lb2fft)

Please read the file above carefully, and make sure your system works if the commands above are used.

In the end-points listed above, whenever we have not specified a certain form of validity checking on the given input, then you can assume that the input will be valid in that way and you need not check for it.

You need not use databases in Phase 1. Each microservice can represent its data using in-memory data structures. When the Restaurant microservice starts up it will initialize the inventory from the /initialData.txt file. The Delivery microservice will initialize all its data also from this file, will initialize the status of each agent mentioned in /initialData.txt to *signed-out* state, and will have an empty set of orders. The Wallet microservice will initialize the balances of all customers in the /initialData.txt file to the given initial balance in this file.

In Phase 1, you can assume that all requests will be sequential (i.e., no concurrent requests to any of the end-points). This will change in Phase 2.

## Test cases

Public Test Case 1:

[https://indianinstituteofscience-my.sharepoint.com/:u:/g/personal/raghavan\\_iisc\\_ac\\_in/ETY6SlYQdhKnTWNsGwYISwBh1\\_aHiH\\_jrLM1POxYZ\\_DYQ?e=XjIWoc](https://indianinstituteofscience-my.sharepoint.com/:u:/g/personal/raghavan_iisc_ac_in/ETY6SlYQdhKnTWNsGwYISwBh1_aHiH_jrLM1POxYZ_DYQ?e=XjIWoc)

Public Test Case 2:

[https://indianinstituteofscience-my.sharepoint.com/:u:/g/personal/raghavan\\_iisc\\_ac\\_in/EXKdmDxr-xlBkfGLAnNsjiMB7cDQcfXYing5bfnSTceJQw?e=eZszz](https://indianinstituteofscience-my.sharepoint.com/:u:/g/personal/raghavan_iisc_ac_in/EXKdmDxr-xlBkfGLAnNsjiMB7cDQcfXYing5bfnSTceJQw?e=eZszz)

In addition to using the two test cases above, you are to develop your own sufficient set of test cases. Each test case should finally print "Pass" or "Fail" exactly once. Note, 'Pass' should mean we got the expected outcome, while 'Fail' should mean we got an unexpected outcome. In some cases, such as Test Case 2 above, a rejected request is the expected outcome. As a good practice, a single test case should not test multiple independent scenarios, as you should prefer to have separate test cases for separate scenarios. It would be a good practice to begin each test case with invocations to all three reinitialize end-points, so that the internal data of each microservice gets re-initialized. **You should aim to create a good set of test cases that test numerous interesting scenarios, corner cases, etc.**

You can implement your test cases using Python, shell scripts (using CURL), Java, or any other suitable language. Each test case should be a standalone program that we can

execute. Place your test cases in a separate folder (separate from the three Spring projects). Include a README file in this folder with instructions on how to run the test cases. Please make sure we don't need to have any uncommon packages installed on our side to run your test cases. If your test cases do need any package, mention the Ubuntu package name to be installed.

## Initialization

An input file (the same input file) will be given to all three of your microservices. This will be mounted as file /initialData.txt using the "-v" option of "docker run" and available inside all your containers. You should treat this file as read-only in your code. A sample contents for this file is as follows:

```
101 2 //Restaurant_Id  number_of_item_available
1 180 10 //itemId price initial_quantity
2 230 20
102 3 //Restaurant_Id  number_of_item_available
1 50 10 //itemId price initial_quantity
3 60 20
4 45 15
****
201 // DeliveryAgent_Id
202
203
****
301 //Customer_Id
302
303
****
2000 //Initial wallet balance of all customers
```

[The comments will be not present in the data file.]

## Logistics

Each project is to be done by a team of students. Discussion across teams is not allowed. If you need any clarifications, please post in the "Project 1 (Spring)" channel on Teams. You can look at generic Java or Spring code fragments on the internet, but you should not look at any code projects online that resemble this given project.

You will implement this entire 1st project in two phases. **The deadline to complete the first phase is Feb. 8th 11.59 pm.** There will be no extensions granted under any circumstance. We will share with you instructions soon on how to upload your projects for evaluation. We will first evaluate your system by running it without your involvement. The correct

performance of your system at this time on the public test cases, our private test cases, and your test cases, with exact expected output in each test case, will account for 75% of the marks of the 1st phase. The remaining 25% marks will be decided on the basis of a brief viva per team, where you will explain your code and the test cases you developed. The elegance and completeness of your code, the comments in your code, and the completeness of your test cases, will be the factors here.

The first phase will have its own marks, which are assigned on the basis of your first phase submission. If the first phase demo shows any incompleteness/incorrectness, then you will not get full marks for this phase; however, you should still fix the first phase's issues after the first phase is over before you begin the work of the second phase as otherwise the 2nd phase performance will also suffer. (Phase 2 code will be based on Phase 1 code.)

## Phase 2 Requirements

- The deadline to complete and upload this phase is March 3rd 11.59pm. The deadline will be absolutely strict.
- The microservices and end-points remain exactly the same as in Phase 1.
- Deploy and expose all three microservices as load-balanced services under minikube.
- For simplicity, fix the number of replicas of Restaurant and Wallet to one replica each. Also, they don't need to use databases. They can continue to keep their data in-memory data structures. (Therefore, these microservices will not be resilient to crashes and will be stateful.)
- To ensure correctness of the Restaurant microservice, ensure that any two end-point executions that update information for the same restaurant execute in isolation even if the requests to these end-points arrive concurrently. That is, they should execute one after the other. Also, the /reInitialize end-point should execute in isolation from all other endpoint executions pertaining to all restaurants. Your code should ensure these isolation properties even if concurrent requests arrive.
- To ensure correctness of the Wallet microservice, ensure that any two end-point executions that update **or access** information for the same wallet execute in isolation. Also, the /reInitialize end-point should execute in isolation from all other endpoint executions pertaining to all wallets. Your code should ensure all these isolation properties even if concurrent requests arrive to the microservice.
- Have a variable number of Delivery replicas, from 1-3, which changes at runtime based on the load. The Delivery replicas should be stateless, so that any request can be directed to any replica safely.
- All Delivery replicas should use a common database service that is deployed as a separate service within the minikube cluster. The Delivery replicas should keep all information that was in-memory in Phase 1 in this database. You can use any kind of database system as you like, and should use Spring Data JPA to access the database. Each execution of each end-point should be treated as a serializable transaction as far as the database is concerned. **The /reInitialize end-point of Delivery should clear the contents of the table(s) and re-initialize them from /initialData.txt.**

- The Delivery service as a whole should be resilient to crashes of the microservice replicas.
- The test cases will be similar to the Phase 1 test cases, but this time ~~time multiple test cases will be executed at the same time concurrently~~ **you need to provide concurrent test cases**. Note, if we take two test cases from Phase 1 and they both used to pass (individually) in Phase 1, they may not pass in Phase 2 when executed concurrently due to interactions between them. Therefore, the assertions **in your concurrent test cases** (i.e., checks on http responses) need to be adjusted to account for possible interactions. Here is a concurrent public test case:  
[https://indianinstituteofscience-my.sharepoint.com/:u:/g/personal/raghavan\\_iisc\\_ac\\_in/ESXKMdZ40kJAultd0k0DuvEBYfgTF7KDmY\\_bcwAJrYyGTQ?e=PnBd9k](https://indianinstituteofscience-my.sharepoint.com/:u:/g/personal/raghavan_iisc_ac_in/ESXKMdZ40kJAultd0k0DuvEBYfgTF7KDmY_bcwAJrYyGTQ?e=PnBd9k)
- The services should internally contact each other via their minikube service names.
- ~~Regarding external visibility of the services, after “minikube tunnel” is run, the three services will have IP addresses that are visible on your development machine. However, using the port forwarding option within minikube, you should further ensure that the three services are available, respectively, at ports 8080, 8081, and 8082, on <http://localhost>.~~
- ~~You will need to provide a single “launch” script that builds your docker images, deploys them, exposes the services, sets up the port forwarding, etc., so that the test cases are ready to run when your script completes. We will provide more information about this script later.~~
- Please provide a `launch.sh` script. This script should perform all the following actions:
  - Do “minikube start”, followed by “eval \$(minikube docker-env)”
  - Compile all your three projects and build their docker images directly into the minikube container. For easy identification, please ensure that each image’s name contains one of your team member’s name followed by the string “restaurant”/“delivery”/“wallet”.
    - A new requirement: Please make your code be Java 11 compatible, and use a Java 11 base image while creating your docker images
  - Launch all four deployments and services (three from your images, the fourth corresponding to the database) within minikube. Follow the same naming convention for the deployments and services as for the images.
  - Use the port forwarding option of minikube to expose your main three services at ports 8080, 8081, and 8082, respectively, on <http://localhost>, so that the test cases can access these services. (The database service need not be exposed to the test cases.)
  - Important note: We will no longer be using the commands mentioned in the following file to test your system:  
[https://indianinstituteofscience-my.sharepoint.com/:t:/g/personal/raghavan\\_iisc\\_ac\\_in/Eee9OQY3WNZKt4EtlvdgeccBvaiT5hMEc5hnbT1RD4545w?e=lb2fft](https://indianinstituteofscience-my.sharepoint.com/:t:/g/personal/raghavan_iisc_ac_in/Eee9OQY3WNZKt4EtlvdgeccBvaiT5hMEc5hnbT1RD4545w?e=lb2fft)  
 Instead, we will be executing your `launch.sh` script first, and then running the test cases one by one.
- Provide a `teardown.sh` file that performs the following actions:
  - Permanently delete all database tables you created



- Delete all your deployments and services permanently from minikube
  - Do “minikube stop”
- As in Phase 1, you will need to submit a folder that contains four sub-folders: three corresponding to your three projects, and one containing your test cases. Your top-level folder should also contain launch.sh and teardown.sh.
- Your test-cases sub-folder should contain concurrent test-cases only. (Of course, the old Phase 1 sequential test cases should continue to pass.) Include a comment at the top of each concurrent test case saying what the test case is testing, intuitively. This sub-folder should also contain a brief README file on how to execute your test-cases, and any non-standard packages we need to install to run your test cases (hopefully none).
- If any bugs were identified in your Phase 1, please fix them in your Phase 2 submission.
- The performance of your system on the public test cases, our private test cases, and your test cases, with exact expected output in each test case, will account for 75% of the marks of the 2nd phase. The remaining 25% marks will be based on the elegance and completeness of your code, the comments in your code in each function and in each test case, and on sufficient coverage of scenarios in your test cases.

#### Resources on Spring Data JPA:

- Geetam’s lecture given on Feb. 5th (recording available)
- A quick tutorial to get started: <https://spring.io/guides/gs/accessing-data-jpa/>
- A more detailed tutorial: <https://www.petrikainulainen.net/spring-data-jpa-tutorial/>